



CS 5330: Pattern Recognition and Computer Vision

Northeastern University

OpenCV Workshop

Lab 5: Histogram

** Contributed by Fall 2024 TAs: Byunghyun Ko, Yihan Wang and Taiwei Cui*

Histogram

1. Introduction to Histograms
2. Types of Histograms
3. Grayscale Histogram Implementation
4. Color Histogram Implementation

Introduction to Histograms

- Histogram: graphical representation that shows the frequency of pixel intensity values in an image
 - Pixel Intensity: ranges from 0 (black) to 255 (white) for grayscale
 - Usage of Histogram: can help understand how pixels are distributed across an image – this is particularly helpful with image thresholding, contrast adjustment, and object detection

Types of Histograms

- Grayscale Histogram: plots the intensity values of pixels in grayscale image
 - The x-axis represents pixel intensity (0-255) and the y-axis represents the number of pixels with each intensity
 - Usages: image brightness and contrast analysis, thresholding, image equalization, image matching and comparison, object detection, etc.
- Color Histogram: created separately for each of the three color channels – tells you how intensities are distributed for each color
 - Usages: image classification and retrieval, color enhancement and correction, color segmentation dominant color extraction, object recognition

Grayscale Histogram Implementation

- Load an image and convert it to grayscale

Code

```
import cv2
import numpy as np
# Load the image in grayscale mode
img = cv2.imread('image.jpg', 0)
```

- “0” in *cv2.imread()* tells OpenCV to load the image in grayscale mode

Grayscale Histogram Implementation

- Calculate the Histogram:

```
hist = cv2.calcHist([img], [0], None, [256], [0, 256])
```

- `cv2.calcHist()`: This OpenCV function computes the histogram of an image.
 - **[img]**: The input image, which is passed as a list because OpenCV expects the image data in this format.
 - **[0]**: The index of the channel for which the histogram is calculated. Since the image is grayscale, there's only one channel (0).
 - **None**: This is where a mask would go if you wanted to calculate the histogram for only a specific region of the image. Here, we use the entire image, so no mask is needed.
 - **[256]**: This represents the number of bins. Since the intensity values of the grayscale image range from 0 to 255, we use 256 bins to represent all possible values.
 - **[0, 256]**: The range of pixel values we want to consider (0 to 255).

Grayscale Histogram Implementation

- Plot the histogram using Matplotlib

Code

```
import matplotlib.pyplot as plt
# Plot the grayscale histogram
plt.plot(hist)
plt.title('Grayscale Histogram')
plt.xlabel('Pixel Intensity')
plt.ylabel('Frequency')
plt.show()
```

- `plt.plot(hist)`: This function plots the computed histogram, showing how often each intensity value occurs in the image. The **x-axis** represents the pixel intensity values (from 0 to 255).
- The **y-axis** represents the frequency, i.e., how many times each intensity value appears in the image.

Color Histogram Implementation

- Load an image in color mode

Code

```
import cv2  
Import numpy as np  
# Load the image in color mode  
img = cv2.imread('image.jpg')
```

- Note the absence of “0” in parameters -> not loading it in grayscale mode

Color Histogram Implementation

Code

```
# Initialize colors for BGR channels
colors = ('b', 'g', 'r')

# Loop through each color channel (Blue, Green, Red)
for i, col in enumerate(colors):
    # Calculate the histogram for each channel
    hist = cv2.calcHist([img], [i], None, [256], [0, 256])

    # Plot the histogram for the current channel
    plt.plot(hist, color=col)

plt.title('Color Histogram')
plt.xlabel('Pixel Intensity')
plt.ylabel('Frequency')
plt.show()
```

Color Channels (BGR): OpenCV uses the BGR (Blue, Green, Red) color model by default. We calculate a separate histogram for each channel.

- `colors = ('b', 'g', 'r')`: This tuple represents the three color channels—Blue, Green, and Red.
- `for i, col in enumerate(colors)`: This loop iterates over each color channel. `i` is the channel index, and `col` is the color name ('b', 'g', or 'r').

• `cv2.calcHist([img], [i], None, [256], [0, 256])`: This computes the histogram for the current channel (`i`). **[img]**: The image from which the histogram is being computed.

- **[i]**: The index of the color channel (0 for Blue, 1 for Green, 2 for Red).
- **None**: We're not using a mask, so all pixels in the channel are considered.
- **[256]**: The number of bins to represent the pixel intensities (0-255).
- **[0, 256]**: The range of pixel intensities to consider for the histogram (0-255).

- `plt.plot(hist, color=col)`: This command plots the histogram for the current channel. The **color** argument ensures that the histogram is plotted in the respective channel's color. For the Blue channel, it plots the histogram in blue.
- For the Green channel, it plots the histogram in green.
- For the Red channel, it plots the histogram in red.

Summary

- Grayscale Histogram:
 - Only one channel, representing the intensity of light from 0 (black) to 255 (white).
 - Used to understand image brightness, contrast, and for basic operations like thresholding.
- Color Histogram:
 - Three separate histograms, one for each color channel (Blue, Green, Red).
 - Helps analyze the color distribution in an image, used for tasks like object recognition, color enhancement, and image classification.